



**Distributed Computer Systems Lab**

<http://disco.informatik.uni-kl.de>



# OpenSky's Data, Algorithms, and Tools

Matthias Schäfer

I'd like to use  
OpenSky's  
data! How does  
it look like??



# Message Format

- Avro files contain serialized sensor data with the following schema:

```
{
  "name": "ModeSEncodedMessage",
  "type": "record",
  "namespace": "org.opensky.avro.v2",
  "fields": [
    {"name": "sensorType", "type": "string"},
    {"name": "sensorLatitude", "type": ["double", "null"]},
    {"name": "sensorLongitude", "type": ["double", "null"]},
    {"name": "sensorAltitude", "type": ["double", "null"]},
    {"name": "timeAtServer", "type": "double"},
    {"name": "timeAtSensor", "type": ["double", "null"]},
    {"name": "timestamp", "type": ["double", "null"]},
    {"name": "rawMessage", "type": "string"},
    {"name": "sensorSerialNumber", "type": "int"},
    {"name": "RSSIPacket", "type": ["double", "null"]},
    {"name": "RSSIPreamble", "type": ["double", "null"]},
    {"name": "SNR", "type": ["double", "null"]},
    {"name": "confidence", "type": ["double", "null"]}
  ]
}
```

# A Typical Instance

```
{  
  'sensorType':          'OpenSky',  
  'sensorSerialNumber':  699700118,  
  'sensorLatitude':      50.04854,  
  'sensorLongitude':     8.4879,  
  'sensorAltitude':     121.0,  
  'timeAtServer':       1441929601.158738,  
  'timestamp':          86844078.0,  
  'rawMessage':         '8d4054a760c380936e7a746ffa90'  
}
```

# Example Message: Sensor Information

```
{  
  'sensorType':          'OpenSky',  
  'sensorSerialNumber':  699700118,  
  'sensorLatitude':      50.04854,  
  'sensorLongitude':     8.4879,  
  'sensorAltitude':      121.0,  
  'timeAtServer':        1441929601.158738,  
  'timestamp':           86844078.0,  
  'rawMessage':          '8d4054a760c380936e7a746ffa90'  
}
```

- `sensorType` **determines** availability of metadata and how to interpret it
  - E.g. accuracy of timestamp, availability of RSSI, ...
- `sensorSerial` **is** the unique identifier of each sensor device
- `sensorLatitude` **and** `sensorLongitude` **in** decimal degrees
- `sensorAltitude` **in** feet

# Example Message: Timestamps

```
{  
  'sensorType':          'OpenSky',  
  'sensorSerialNumber':  699700118,  
  'sensorLatitude':      50.04854,  
  'sensorLongitude':     8.4879,  
  'sensorAltitude':      121.0,  
  'timeAtServer':        1441929601.158738,  
  'timestamp':           86844078.0,  
  'rawMessage':          '8d4054a760c380936e7a746ffa90'  
}
```

- `timeAtServer` is unix timestamp for the arrival at the OpenSky server
  - Contains Internet jitter
- `timestamp` is a hardware timestamp for the message detection
  - Unit and accuracy depend on hardware
  - Can be used e.g. for multilateration

# Example Message: Timestamps

```
{  
  'sensorType':          'OpenSky',  
  'sensorSerialNumber':  699700118,  
  'sensorLatitude':      50.04854,  
  'sensorLongitude':     8.4879,  
  'sensorAltitude':      121.0,  
  'timeAtServer':        1441929601.158738,  
  'timestamp':           86844078.0,  
  'rawMessage':          '8d4054a760c380936e7a746ffa90'  
}
```

- rawMessage is the Mode message in hex representation
  - CRC is set to 000000 from some sensors
- A complete open-source ADS-B decoder is avlbl on our github account
  - <https://github.com/openskynetwork/java-adsb>

Mhmm... okay,  
and what to do  
when I get a pile  
of data?





# Example file:



raw2015091100.avro  
(Size: 8.5GB)

Tools are available in decoder.jar at  
[github/openskynetwork/osky-sample](https://github.com/openskynetwork/osky-sample)


# Tool #1: AvroInfo

- Provides highlevel information about OpenSky's avro files


```
$ java -cp decoder.jar org.opensky.tools.AvroInfo raw2015091100.avro

...
Counting entries: 94918705
Earliest entry: Fri Sep 11 02:00:00 CEST 2015
Latest entry: Sat Sep 12 01:59:59 CEST 2015
...
```

Messages/Sensor:




Message Type Distribution:



## Tool #2: Avro2KML

- Want to get a visual idea of the data? Let's generate a KML file containing the first 1000 flights of our sample.
- Keyhole Markup Language (KML) used to visualize data in Google Earth


```
$ java -cp decoder.jar org.opensky.tools.Avro2Kml -0 -n 1000  
raw2015091100.avro 1000.kml  
  
...  
2015-11-24 12:53:09 WARN PositionDecoder - Position messages  
should be ordered!  
2015-11-24 12:53:09 WARN PositionDecoder - Position messages  
should be ordered!  
2015-11-24 12:53:09 WARN PositionDecoder - Position messages  
should be ordered!  
...
```



Oops... what happened?

# Primer on Position Decoding in ADS-B

- ADS-B uses Compact Position Reporting (CPR)
  - Encoding saves 10 bits per airborne position (14 bits per surface position)
- CPR uses special coordinate system for encoding/decoding
  - There are even and odd positions
  - **Each of them alone cannot be decoded uniquely**
- Each ADS-B position report contains either even or odd position, not both ☹️
- ADS-B position decoding:  
To get unique position we need an even and an odd message (global CPR) or a nearby reference position (local CPR)
  - For global CPR being applicable, both reports (even and odd) must encode positions close to each other



red lines: "even zones"  
blue lines: "odd zones"

**Ah, I see! That's why  
position messages have  
to be decoded in order!  
But how can I sort the  
avro file?**




## Tool #3: AvroSort

- Sorts all entries of an avro file by their `timeAtServer`

```
$ java -cp decoder.jar org.opensky.tools.AvroSort  
raw2015091100.avro raw2015091100_sorted.avro
```

```
Warning: make sure you have enough main memory. Otherwise, use  
AvroSplit first.
```



- But be careful: Sorting the file requires keeping the whole data into RAM
  - The file has 8.5 Gbyte! Loading it completely to RAM consumes even more space.
- Idea: Split file before sorting

# Tool #4: AvroSplit

- Splits avro files into  $n$  smaller pieces with the following properties:
  - The pieces have roughly the same size
  - All messages of messages of an aircraft are in the same file (important for position decoding!)

```
$ java -cp decoder.jar org.opensky.tools.AvroSplit -n 10
raw2015091100.avro -o raw2015091100_part

...
$ ls -alh raw2015091100_part*
-rw-r--r-- 1 matze users 788M Nov 25 16:26 raw2015091100_part1.avro
-rw-r--r-- 1 matze users 899M Nov 25 16:26 raw2015091100_part10.avro
-rw-r--r-- 1 matze users 842M Nov 25 16:26 raw2015091100_part2.avro
-rw-r--r-- 1 matze users 859M Nov 25 16:26 raw2015091100_part3.avro
-rw-r--r-- 1 matze users 988M Nov 25 16:26 raw2015091100_part4.avro
-rw-r--r-- 1 matze users 791M Nov 25 16:26 raw2015091100_part5.avro
-rw-r--r-- 1 matze users 898M Nov 25 16:26 raw2015091100_part6.avro
-rw-r--r-- 1 matze users 854M Nov 25 16:26 raw2015091100_part7.avro
-rw-r--r-- 1 matze users 834M Nov 25 16:26 raw2015091100_part8.avro
-rw-r--r-- 1 matze users 872M Nov 25 16:26 raw2015091100_part9.avro
```

# Sorting continued...

- Now we can sort the smaller files

```
$ for f in raw2015091100_part*; do java -cp decoder.jar  
  org.opensky.tools.AvroSort ${f} sorted_${f}; done  
...  
$ ls sorted_raw2015091100_part*  
sorted_raw2015091100_part1.avro  sorted_raw2015091100_part5.avro  
sorted_raw2015091100_part10.avro sorted_raw2015091100_part6.avro  
sorted_raw2015091100_part2.avro  sorted_raw2015091100_part7.avro  
sorted_raw2015091100_part3.avro  sorted_raw2015091100_part8.avro  
sorted_raw2015091100_part4.avro  sorted_raw2015091100_part9.avro
```



But I would like to  
decode the whole  
dataset! Not only  
a 10th at a time...



# Joining + Generating KML Files

- AvroSplit also accepts multiple input files

- It can be used to concatenate files by setting `n=1`

```
$ java -cp decoder.jar org.opensky.tools.AvroSplit -n 1  
sorted_* -o raw2015091100_sorted
```

- Note: The resulting file (`raw2015091100_sorted1.avro`) is not completely ordered. Only the messages of a certain aircraft are ordered by `timeAtServer`!

- And now try generating the KML again:

```
$ java -cp decoder.jar org.opensky.tools.Avro2Kml -0 -n 1000  
raw2015091100_sorted1.avro 1000.kml
```

- Have a look at `1000.kml` with Google Earth ☺

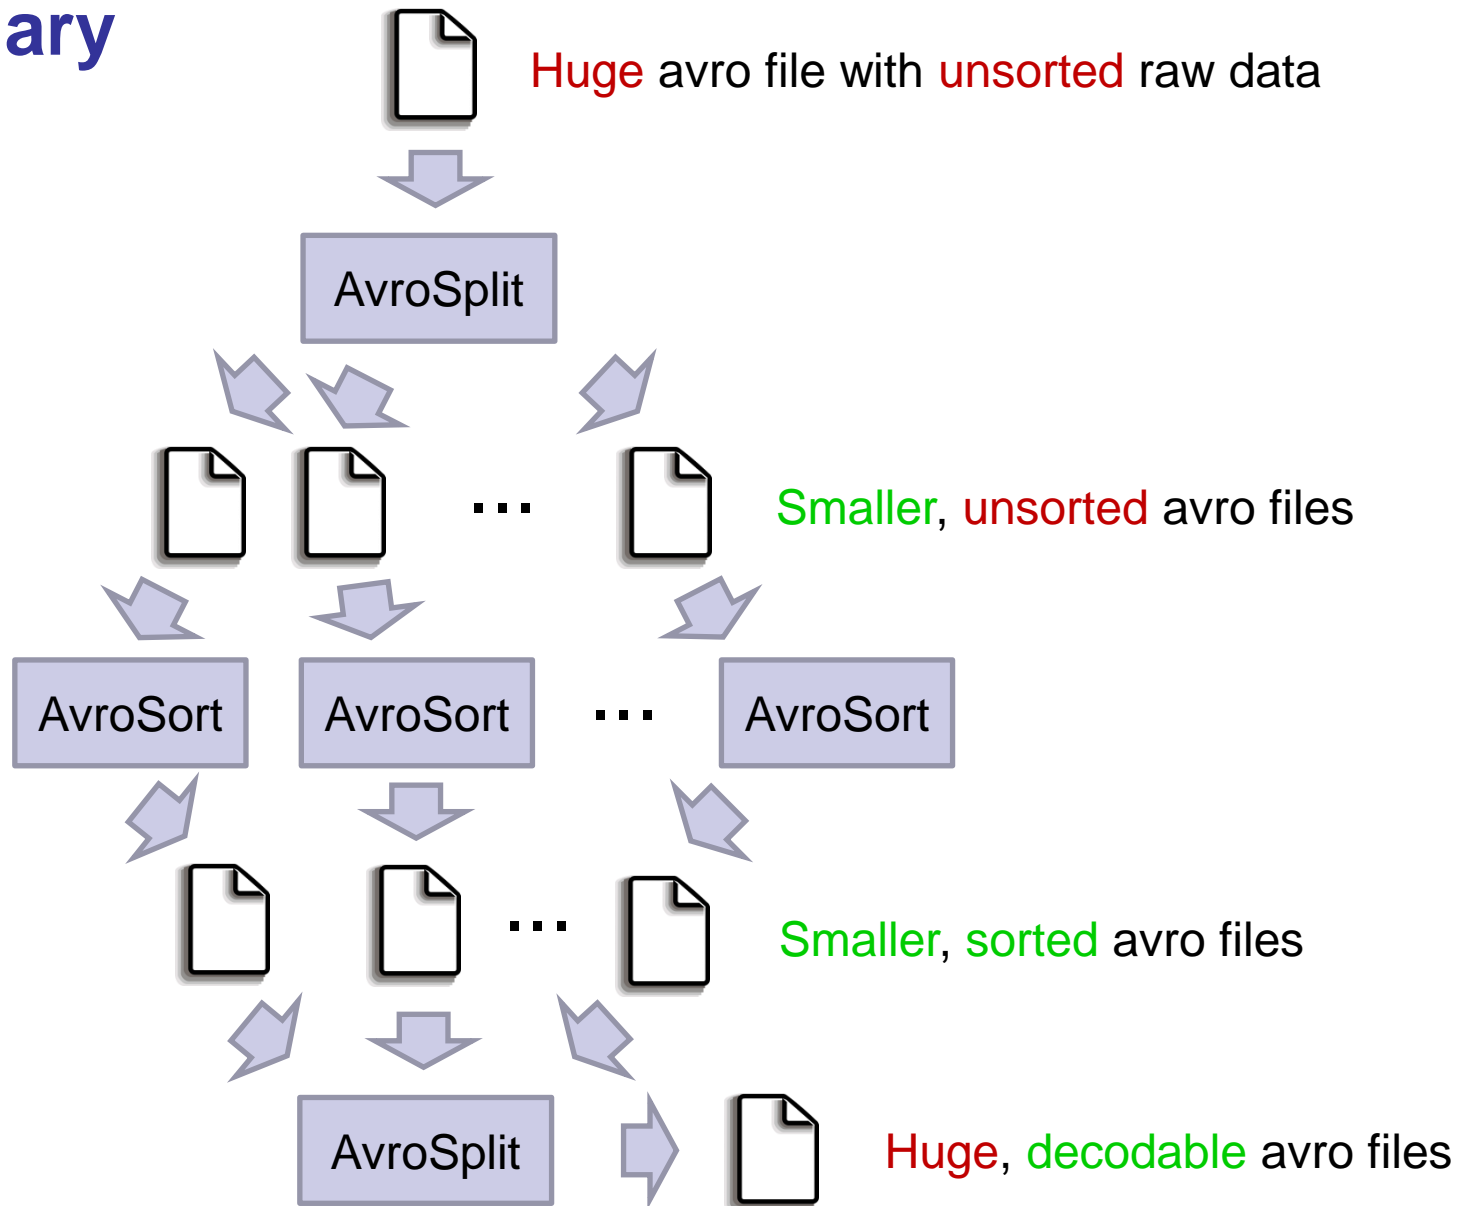








# Summary



I want to have  
some statistics  
and analyse flight  
paths!!



## Tool #5: Avro2SQLite

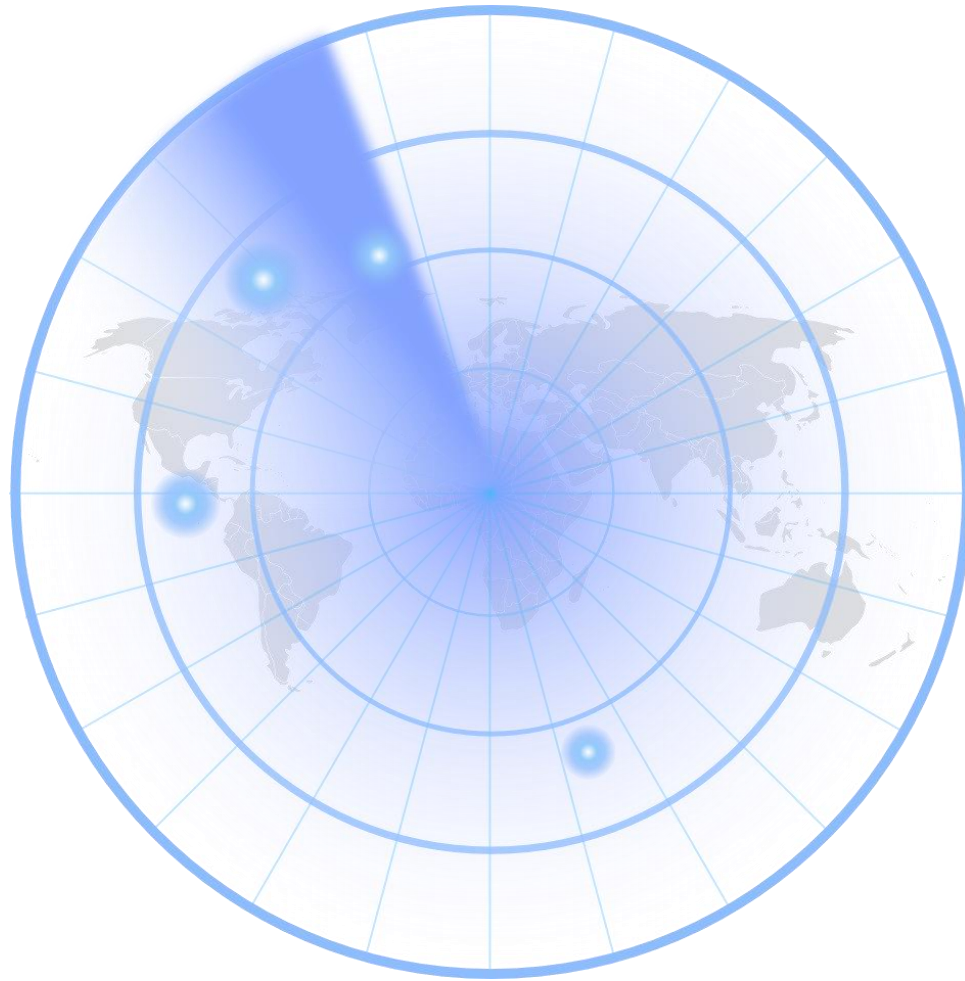
- Generates SQLite database with decoded flights including their positions and velocities
  - Can be queried and loaded for your analyses
  - But be careful: it does not scale ☹

```
$ java -cp decoder.jar org.opensky.tools.Avro2SQLite
raw2015091100_sorted1.avro raw2015091100.sqlite
...
$ sqlite3 raw2015091100.sqlite
sqlite> SELECT * FROM flights ORDER BY RANDOM() LIMIT 1;
54951|1441986565.60435|1441989692.76459|a2b728|UPS213
sqlite> SELECT * FROM positions WHERE flight=54951;
54951|1441986828.41693|-3.00107247488836|53.4179077148437|10668.0|370.4
54951|1441986829.43119|-2.99722726004461|53.4175415039062|10668.0|370.4
...
sqlite> SELECT * FROM velocities WHERE flight=54951;
54951|1441986606.30583|253.06671140109|97.240590129645|0.0
54951|1441986611.21757|253.387563099794|96.8797324654071|0.65024
...
```

Feel free to extend our tools or create new ones! All source codes and examples can be found on github:

<https://github.com/openskynetwork/java-adsb>





Thanks 😊